# YESWANTH ARASAVALLI

Data Analyst | Data Scientist

Hyderabad, India • +91 6303321322 • maneswararasavalli@gmail.com • linkedin.com/in/yeswantharasavalli • yeswantharasavalli.me

## SUMMARY

M.S. in Computer Science (Big Data Analytics) with three data-focused internships spanning SQL, ETL, Power BI/Tableau dashboarding, and applied machine learning. Comfortable across the full analytics workflow — data cleaning, SQL, statistical analysis, model building and evaluation, and dashboard delivery — with additional independent project experience in NLP, computer vision, and time-series forecasting. Currently doing freelance, project-based AI data-annotation and prompt-evaluation work.

## TECHNICAL SKILLS

**Programming & SQL:** Python, SQL (CTEs, window functions, query optimisation), R

**Data Analytics & BI:** Excel, Power BI, Tableau, KPI analysis, dashboard development, exploratory data analysis, statistical analysis, hypothesis testing

**Data Processing & Engineering:** Pandas, NumPy, data cleaning, data transformation, ETL pipelines, feature engineering, PostgreSQL, MongoDB

**Machine Learning:** Scikit-learn, XGBoost, Random Forest, Logistic Regression, model evaluation, classification, segmentation

**Deep Learning & NLP:** PyTorch, LSTM, DistilBERT, Hugging Face Transformers, TF-IDF, NLP

**Computer Vision:** EfficientNet-B0, image classification

**Tools / Deployment (Project-Level):** FastAPI, REST APIs, Docker, MLflow, Evidently, Streamlit, Gradio, Ollama, Git

**Cloud:** AWS, GCP, Azure — project/certification exposure

## EXPERIENCE

**AI Data Annotation Specialist & Prompt Engineer — Freelance / Project-Based** Feb 2024 – Present

*TELUS International — Remote*

- Performed QA review on NLP and CV annotation batches, resolved ambiguous edge cases, and documented labelling rules to reduce inconsistency across review cycles.
- Evaluated prompt variants for generative AI tasks and tracked output quality across iterations to identify recurring failure patterns.
- Flagged systematic labelling errors and proposed structural fixes to improve training-data quality before model ingestion.

**Data Science Intern — Postulate IRT, India** Jan 2024 – Feb 2024

- Cleaned and standardised 150K+ NLP speech records in Python, resolving missing values, encoding issues, and schema mismatches.
- Built and benchmarked Ridge, XGBoost, and ARIMA forecasting models; XGBoost outperformed the baseline by 22% on the primary validation metric.
- Built Tableau dashboards with calculated KPI fields and drill-through filters for weekly reporting.

**Data Analyst Intern — Edunet Foundation, India** Jul 2023 – Jan 2024

- Wrote SQL pipelines using CTEs and window functions to consolidate data from cloud and on-prem sources, reducing query runtime by 25%.
- Developed customer segmentation models using tree-based classifiers, evaluated on precision/recall/F1, and delivered segment definitions to the business team.
- Built Power BI dashboards tracking operational KPIs; worked with stakeholders to redefine two metrics.

**Data Science Intern — PHN Technology Pvt. Ltd., India** Apr 2023 – Jul 2023

- Automated ETL scripts for log data ingestion using Python and SQL, replacing 3–4 hours/week of manual exports and improving pipeline reliability.
- Built Tableau and Excel reports that cut manual reporting effort by 40% by standardising data pulls and templating recurring layouts.
- Built and deployed R-based statistical models through IBM Watson Studio to support repeatable internal analytics workflows.

## INDEPENDENT PROJECTS

**Domain-Agnostic Decision Intelligence System** | *Python, Scikit-learn, Pandas, Streamlit, Flask*

- Built a multi-stage pipeline (data quality, domain auto-detection across 13 domains, KPI computation, driver-graph analysis, scenario simulation) producing a 10-section report.
- Added a Flask REST API layer (7 endpoints); validated with 454 test functions across 30 test files.

**AI Job Application Intelligence System** | *Python, Ollama LLM, Streamlit, SQLite, FastAPI, Playwright*

- Built an end-to-end local system: JD processing, LLM-based resume tailoring, 6-dimension ATS scoring, and an APPLY/MAYBE/IMPROVE decision engine, with an 8-page Streamlit dashboard and a Playwright-based LinkedIn scraper.
- Covered by 144 automated tests across 23 test files; zero cloud dependencies.

**Amazon Food Reviews Sentiment Analysis** | *Python, DistilBERT, PyTorch, Hugging Face, Gradio*

- Fine-tuned DistilBERT on 568K+ reviews, reaching ~94% accuracy vs. ~89% TF-IDF + Logistic Regression baseline; evaluated with F1, ROC, and confusion matrix.
- Deployed a Gradio demo on Hugging Face Spaces.

**Retail Product Image Recognition System** | *Python, EfficientNet-B0, PyTorch, MLflow, Streamlit*

- Fine-tuned EfficientNet-B0 on ~200 SKU categories, reaching 84.9% Top-1 accuracy vs. a 76% ResNet-18 baseline; tracked experiments with MLflow.

**Predictive Maintenance — Machine Failure Forecasting** | *Python, PyTorch, LSTM*

- Trained an LSTM on multivariate sensor data for 2-week failure forecasting, reaching 87% validation accuracy and outperforming ARIMA/gradient-boosted baselines on failure recall.

## EDUCATION

**M.S. Computer Science — Big Data Analytics** 2021 – 2023
*Central University of Rajasthan, Ajmer | GPA: 7.3 / 10*

Capstone: Telecom churn prediction using an XGBoost + Random Forest ensemble — 88% accuracy on 200K+ records.

**B.S. Computer Science** 2018 – 2021
*Adikavi Nannaya University, Rajahmundry | GPA: 8.37 / 10*

## CERTIFICATIONS

- Microsoft Certified: Fabric Analytics Engineer Associate — Microsoft
- Google Advanced Data Analytics — Coursera
- Data Science Specialisation — Johns Hopkins University, Coursera
- Machine Learning / AI Engineer — Skillsoft & Infosys Springboard
- MongoDB Python Developer Associate — MongoDB University
- Microsoft Power BI — Infosys Springboard